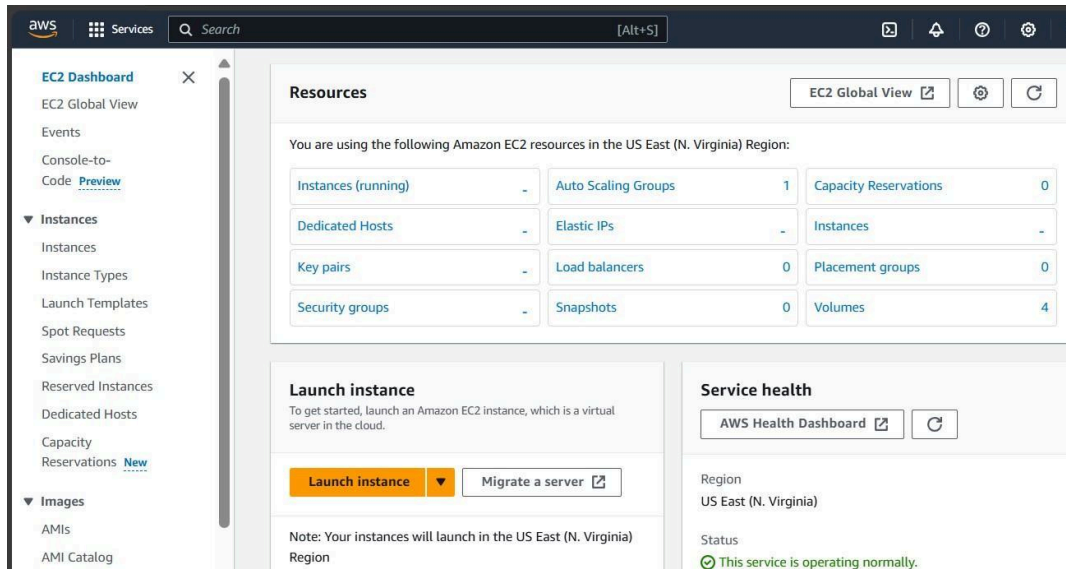


EXPERIMENT NO. 3

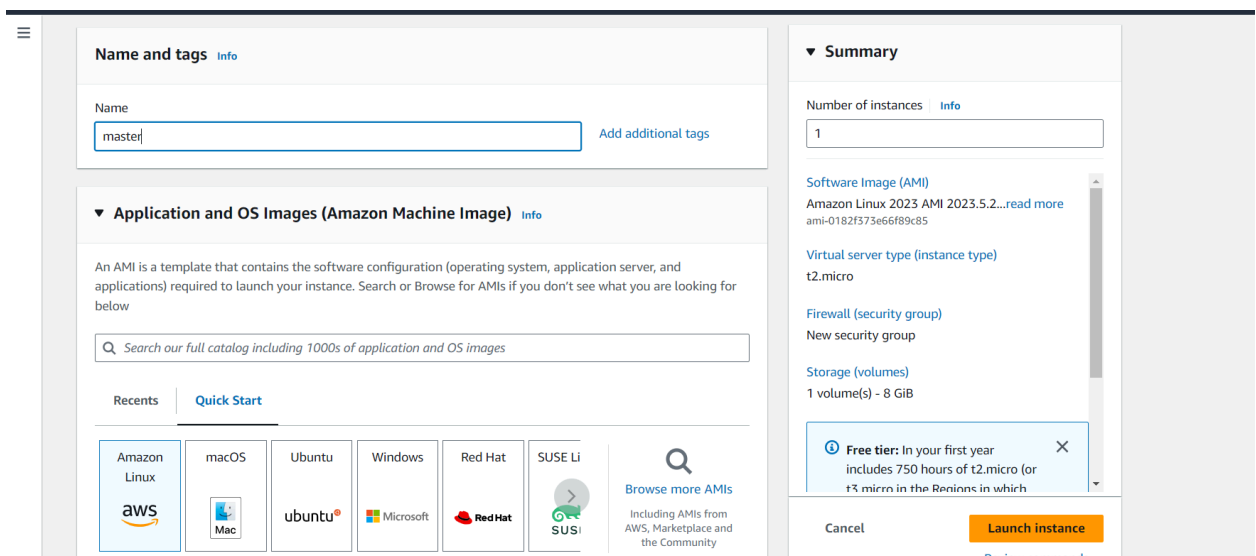
Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud.

Procedure:

1. Creation Of Instance



Search EC-2 instance. Then create three EC-2 instances and choose Amazon Linux as OS and also allow ssh traffic from anywhere.



Description

Amazon Linux 2023 is a modern, general purpose Linux-based OS that comes with 5 years of long term s

▼ Network settings

Info

Edit

Network

Info

vpc-051bba342b3626898

Subnet

Info

No preference (Default subnet in any availability zone)

Auto-assign public IP

Info

Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups)

Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Create security group

Select existing security group

We'll create a new security group called 'launch-wizard-31' with the following rules:

☒ Allow SSH traffic from

Helps you connect to your instance

Anywhere

0.0.0.0/0

☒ Allow HTTPS traffic from the internet

To set up an endpoint, for example when creating a web server

☒ Allow HTTP traffic from the internet

To set up an endpoint, for example when creating a web server

⚠ Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

×

▼ Configure storage

Info

Advanced

1x

8

GiB

gp3

Root volume (Not encrypted)

ⓘ Free tier eligible customers can get up to 30 GiB of EBS General Purpose (SSD) or Magnetic storage

×

▼ Summary

Number of instances

Info

1

Software Image (AMI)

Amazon Linux 2023 AMI 2023.5.2...read more

ami-0182f373e66f89c85

Virtual server type (instance type)

t3.medium

Firewall (security group)

New security group

Storage (volumes)

1 volume(s) - 8 GiB

ⓘ Free tier: In your first year includes

750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million IOs, 1 GiB of snapshots, and 100 GB of bandwidth to the internet.

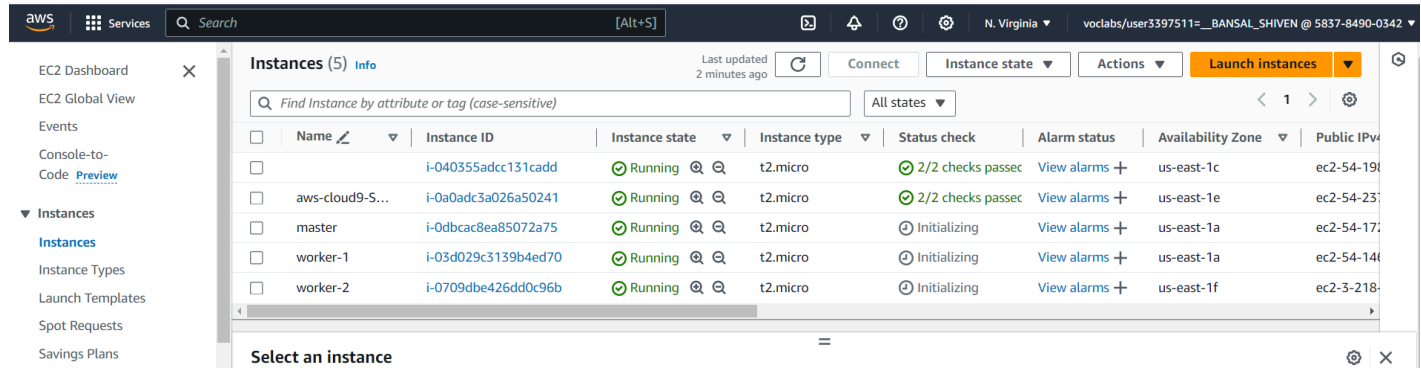
×

Cancel

Launch instance

Review commands

To efficiently run kubernetes cluster select instance type of at least t3.medium as kubernetes recommends at least 2 vCPU to run smoothly on it.



- Then for making connection through SSH into all 3 machines each in separate terminal Use this following command:
`ssh -i <keyname>.pem ubuntu@<public_ip_address>` where keyname is name of the key you created here i created key server.pem and use public IP address.(I have entered this command on git bash where i entered in downloads where server.pem is stored then as the key is not accessible hence we need to change its mode using `chmod 400 "key name.pem"`. Then use the given command for making connections).

```
ADMIN@DESKTOP-LPV2RP5 MINGW64 ~
$ cd Downloads/
```

```
ADMIN@DESKTOP-LPV2RP5 MINGW64 ~/Downloads
$ chmod 400 "server.pem"
```

```
ADMIN@DESKTOP-LPV2RP5 MINGW64 ~/Downloads
$ ssh -i "server.pem" ec2-user@ec2-44-192-73-91.compute-1.amazonaws.com
The authenticity of host 'ec2-44-192-73-91.compute-1.amazonaws.com (44.192.73.91)' can't be established.
ED25519 key fingerprint is SHA256:xIIIIRCBX1LU0Jsm9WHO/YIrDSioPOJfOPEoRKtYJk.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-44-192-73-91.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
```

```

#_
~\##### Amazon Linux 2023
~\#####\
~\###|
~\#/
~\V'-'> https://aws.amazon.com/linux/amazon-linux-2023
~\./
~\./
~\./
[ec2-user@ip-172-31-78-148 ~]$
```

2. Installation Of Docker on three machines

- For installation of Docker into all three machines run the following command:
sudo yum install docker -y

```

Last login: Sat Sep 14 06:53:06 EDT 2024 from 10.0.0.1
[ec2-user@ip-172-31-81-216 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:20:45 ago on Sat Sep 14 06:53:06 2024.
Dependencies resolved.
=====
Package                Arch      Version                               Repository      Size
=====
Installing:
docker                  x86_64    25.0.6-1.amzn2023.0.2               amazonlinux     44 M
Installing dependencies:
containerd              x86_64    1.7.20-1.amzn2023.0.1               amazonlinux     35 M
iptables-libs           x86_64    1.8.8-3.amzn2023.0.2               amazonlinux     401 k
iptables-nft            x86_64    1.8.8-3.amzn2023.0.2               amazonlinux     183 k
libcgroup               x86_64    3.0-1.amzn2023.0.1                 amazonlinux     75 k
libnetfilter_conntrack x86_64    1.0.8-2.amzn2023.0.2               amazonlinux     58 k
libnftnl                 x86_64    1.0.1-19.amzn2023.0.2             amazonlinux     30 k
libnftnl                 x86_64    1.2.2-2.amzn2023.0.2             amazonlinux     84 k
pigz                    x86_64    2.5-1.amzn2023.0.3                 amazonlinux     83 k
runc                    x86_64    1.1.13-1.amzn2023.0.1             amazonlinux     3.2 M
=====

Transaction Summary
=====
Install 10 Packages

Total download size: 84 M
Installed size: 317 M
Downloading Packages:
(1/10): iptables-libs-1.8.8-3.amzn2023.0.2.x86_  5.7 MB/s | 401 kB    00:00
(2/10): iptables-nft-1.8.8-3.amzn2023.0.2.x86_6  6.2 MB/s | 183 kB    00:00
(3/10): libcgroup-3.0-1.amzn2023.0.1.x86_64.rp  2.2 MB/s | 75 kB    00:00
(4/10): libnetfilter_conntrack-1.0.8-2.amzn2023  1.9 MB/s | 58 kB    00:00
(5/10): libnftnl-1.0.1-19.amzn2023.0.2.x86_  1.2 MB/s | 30 kB    00:00
(6/10): libnftnl-1.2.2-2.amzn2023.0.2.x86_64.rp  2.2 MB/s | 84 kB    00:00
(7/10): pigz-2.5-1.amzn2023.0.3.x86_64.rpm      1.7 MB/s | 83 kB    00:00
(8/10): runc-1.1.13-1.amzn2023.0.1.x86_64.rpm    22 MB/s | 3.2 MB   00:00
(9/10): docker-25.0.6-1.amzn2023.0.2.x86_64.rp  32 MB/s | 44 MB    00:01
(10/10): containerd-1.7.20-1.amzn2023.0.1.x86_6  23 MB/s | 35 MB    00:01
-----
Total                                          53 MB/s | 84 MB    00:01
Running transaction check
Transaction check succeeded.

Installing :
Installing : containerd-1.7.20-1.amzn2023.0.1.x86_64
Installing : docker-25.0.6-1.amzn2023.0.2.x86_64
Installing : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
Installing : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
Installing : libcgroup-3.0-1.amzn2023.0.1.x86_64
Installing : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
Installing : libnftnl-1.0.1-19.amzn2023.0.2.x86_64
Installing : libnftnl-1.2.2-2.amzn2023.0.2.x86_64
Installing : pigz-2.5-1.amzn2023.0.3.x86_64
Installing : runc-1.1.13-1.amzn2023.0.1.x86_64

Installed:
containerd-1.7.20-1.amzn2023.0.1.x86_64
docker-25.0.6-1.amzn2023.0.2.x86_64
iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
libcgroup-3.0-1.amzn2023.0.1.x86_64
libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
libnftnl-1.0.1-19.amzn2023.0.2.x86_64
libnftnl-1.2.2-2.amzn2023.0.2.x86_64
pigz-2.5-1.amzn2023.0.3.x86_64
runc-1.1.13-1.amzn2023.0.1.x86_64

Complete!
[ec2-user@ip-172-31-81-216 ~]$ client_loop: send disconnect: Connection reset by
peer

```

- Then, configure cgroup in a daemon.json file by using following commands
cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts":
["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
}

```

    },
    "storage-driver": "overlay2"
  }
}
EOF

```

```

-----
[ec2-user@ip-172-31-81-216 ~]$ cd /etc/docker
[ec2-user@ip-172-31-81-216 docker]$ cat <<EOF | sudo tee /etc/docker/daemon.json

{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
[ec2-user@ip-172-31-81-216 docker]$

```

- Then after this run the following command to enable and start docker and also to load the daemon.json file.

```

sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker

```

```

[ec2-user@ip-172-31-81-216 docker]$ sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-81-216 docker]$ sudo systemctl daemon-reload
[ec2-user@ip-172-31-81-216 docker]$ sudo systemctl restart docker

```

- Then check the version of docker installed. docker -v

```

[ec2-user@ip-172-31-81-216 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc

```

3. Installation Of Kubernetes on three machines

- SELinux needs to be disabled before configuring kubelet thus run the following command
sudo setenforce 0

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

```
[ec2-user@ip-172-31-81-216 ~]$ sudo setenforce 0
[ec2-user@ip-172-31-81-216 ~]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

- Here We are adding kubernetes using the repository whose command is given below. cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo

```
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
```

```
[ec2-user@ip-172-31-81-216 ~]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[ec2-user@ip-172-31-81-216 ~]$
```

- After that Run following command to make the update and also to install kubelet, kubeadm, kubectl: sudo yum update

```
[ec2-user@ip-172-31-81-216 ~]$ sudo yum update
Kubernetes                               114 kB/s | 17 kB      00:00
Dependencies resolved.
Nothing to do.
Complete!
```

```
sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
```

```
Complete!
[ec2-user@ip-172-31-81-216 ~]$ sudo yum install -y kubelet kubeadm kubectl --dis
ableexcludes=kubernetes
Last metadata expiration check: 0:00:35 ago on Sat Sep 14 07:38:02 2024.
Dependencies resolved.
```

Package	Arch	Version	Repository	Size
Installing:				
kubeadm	x86_64	1.30.5-150500.1.1	kubernetes	10 M
kubectl	x86_64	1.30.5-150500.1.1	kubernetes	10 M
kubelet	x86_64	1.30.5-150500.1.1	kubernetes	17 M
Installing dependencies:				
conntrack-tools	x86_64	1.4.6-2.amzn2023.0.2	amazonlinux	208 k
cri-tools	x86_64	1.30.1-150500.1.1	kubernetes	8.6 M
kubernetes-cni	x86_64	1.4.0-150500.1.1	kubernetes	6.7 M
libnetfilter_cthelper	x86_64	1.0.0-21.amzn2023.0.2	amazonlinux	24 k
libnetfilter_cttimeout	x86_64	1.0.0-19.amzn2023.0.2	amazonlinux	24 k
libnetfilter_queue	x86_64	1.0.5-2.amzn2023.0.2	amazonlinux	30 k

Transaction Summary

```
Install 9 Packages
```

```
Total download size: 53 M
```

```
Installed size: 292 M
```

```
Downloading Packages:
```

```
(1/9): libnetfilter_cthelper-1.0.0-21.amzn2023. 436 kB/s | 24 kB    00:00
(2/9): libnetfilter_queue-1.0.5-2.amzn2023.0.2. 1.5 MB/s | 30 kB    00:00
(3/9): libnetfilter_cttimeout-1.0.0-19.amzn2023 309 kB/s | 24 kB    00:00
(4/9): conntrack-tools-1.4.6-2.amzn2023.0.2.x86 2.3 MB/s | 208 kB    00:00
(5/9): cri-tools-1.30.1-150500.1.1.x86_64.rpm   34 MB/s | 8.6 MB    00:00
(6/9): kubectl-1.30.5-150500.1.1.x86_64.rpm    21 MB/s | 10 MB    00:00
(7/9): kubeadm-1.30.5-150500.1.1.x86_64.rpm    17 MB/s | 10 MB    00:00
(8/9): kubelet-1.30.5-150500.1.1.x86_64.rpm    32 MB/s | 17 MB    00:00
(9/9): kubernetes-cni-1.4.0-150500.1.1.x86_64.r 21 MB/s | 6.7 MB    00:00
```

```
Verifying      : kubeadm-1.30.5-150500.1.1.x86_64      8/9
Verifying      : kubectl-1.30.5-150500.1.1.x86_64      7/9
Verifying      : kubelet-1.30.5-150500.1.1.x86_64      8/9
Verifying      : kubernetes-cni-1.4.0-150500.1.1.x86_64 9/9
```

```
Installed:
```

```
conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
cri-tools-1.30.1-150500.1.1.x86_64
kubeadm-1.30.5-150500.1.1.x86_64
kubectl-1.30.5-150500.1.1.x86_64
kubelet-1.30.5-150500.1.1.x86_64
kubernetes-cni-1.4.0-150500.1.1.x86_64
libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64
```

```
Complete!
```

```
[ec2-user@ip-172-31-81-216 ~]$ |
```

- After installing Kubernetes, we need to configure internet options to allow bridging.
 - sudo swapoff -a
 - echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
 - sudo sysctl -p

```
-----
[ec2-user@ip-172-31-81-216 ~]$ sudo swapoff -a
[ec2-user@ip-172-31-81-216 ~]$ echo "net.bridge.bridge-nf-call-iptables=1" | sud
o tee -a /etc/sysctl.conf
net.bridge.bridge-nf-call-iptables=1
[ec2-user@ip-172-31-81-216 ~]$ sudo sysctl -p
net.bridge.brdaef-nf-call-iptables = 1
```


4. Perform this ONLY on the Master machine

- Initialize kubernetes by typing below command

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all
[ec2-user@ip-172-31-81-216 ~]$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
--ignore-preflight-errors=all
I0914 07:46:39.398298 28873 version.go:256] remote version is much newer: v1.3
1.0; falling back to: stable-1.30
[init] Using Kubernetes version: v1.30.4
[preflight] Running pre-flight checks
[WARNING NumCPU]: the number of available CPUs 1 is less than the require
d 2
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.81.216:6443 --token f35qak.rrjyqp89bqsomqrn \
--discovery-token-ca-cert-hash sha256:593693973f6b40e7bc61dec2c73c617ba2
26caafee64bc2776ee360c42b6f29c
```

- So after initialization you will get token at the end for joining master and worker. Like here I got this :(save this token as it is required later. Then you can join any number of worker nodes by running the following on each as root.)

```
kubeadm join 172.31.81.216:6443 --token f35qak.rrjyqp89bqsomqrn \
--discovery-token-ca-cert-hash
sha256:593693973f6b40e7bc61dec2c73c617ba226caafee64bc2776ee360c42b6f29c
```

- Also, Copy the mkdir and chown commands from the top and execute them
 mkdir -p \$HOME/.kube
 sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config
 sudo chown \$(id -u):\$(id -g) \$HOME/.kube/config

```
[ec2-user@ip-172-31-81-216 ~]$ mkdir -p $HOME/.kube
[ec2-user@ip-172-31-81-216 ~]$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube
/config
[ec2-user@ip-172-31-81-216 ~]$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Then, add a common networking plugin called flannel file as mentioned in the code.
 kubectl apply -f

<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>


```
[ec2-user@ip-172-31-81-216 ~]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
error: error validating "https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml": error validating data: failed to download open
api: Get "https://172.31.81.216:6443/openapi/v2?timeout=32s": dial tcp 172.31.81.216:6443: connect: connection refused; if you choose to ignore these errors, tu
rn validation off with --validate=false
```

This step gives an error

Conclusion:

In this experiment, we successfully set up a Kubernetes cluster across three Amazon Linux EC2 instances, each equipped with Kubernetes components. The master node was initialized using kubeadm, and pod networking was configured with the Flannel network plugin. Worker nodes were integrated into the cluster through the join command generated during the master node's initialization. The process provided a comprehensive understanding of Kubernetes cluster setup on EC2. However, there was a noticeable delay in the worker nodes connecting to the master node, which may have been caused by network connectivity issues or configuration discrepancies.