

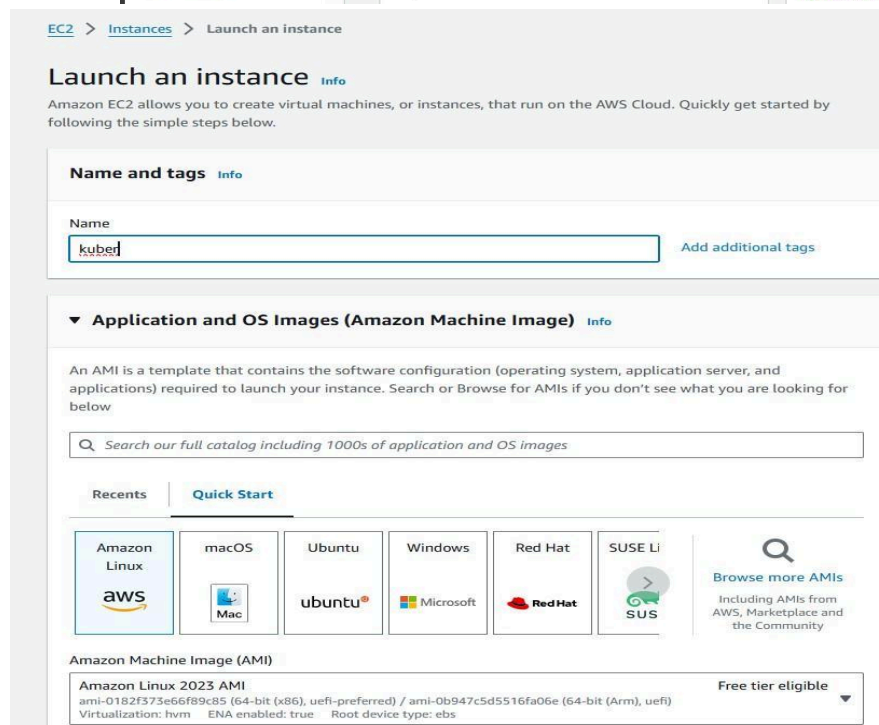
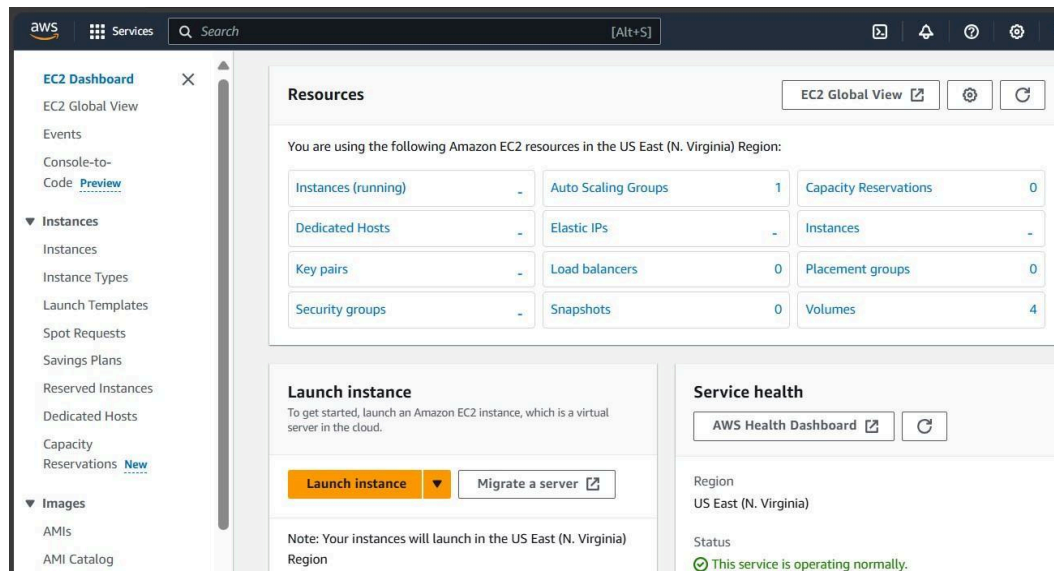
Experiment 4

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

Procedure:

1. Creation Of EC-2 instance

- Create an EC2 AWS Linux instance on AWS .also edit the Security Group Inbound Rules to allow SSH. then select the t2.micro instance type



Network settings [Info](#) [Edit](#)

Network [Info](#)
vpc-0381e49e607677b63

Subnet [Info](#)
No preference (Default subnet in any availability zone)

Auto-assign public IP [Info](#)
Enable
[Additional charges apply](#) when outside of [free tier allowance](#)

Firewall (security groups) [Info](#)
A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Create security group ☐ Select existing security group

We'll create a new security group called 'launch-wizard-8' with the following rules:

- ☒ Allow SSH traffic from
- ☒ Allow HTTPS traffic from the internet
To set up an endpoint, for example when creating a web server
- ☒ Allow HTTP traffic from the internet
To set up an endpoint, for example when creating a web server

Summary

Number of instances [Info](#)

Software Image (AMI)
Amazon Linux 2023 AMI 2023.5.2...[read more](#)
ami-0182f373e66f89c85

Virtual server type (instance type)
t2.micro

Firewall (security group)
New security group

Storage (volumes)
1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which)

[Cancel](#) [Launch instance](#) [Review commands](#)

Instances (6) [Info](#) Last updated less than a minute ago [Refresh](#) [Connect](#) [Instance state](#) [Actions](#) [Launch instances](#)

[All states](#) [1](#) [Settings](#)

<input type="checkbox"/>	Name ↗	Instance ID	Instance state ↗	Instance type ↗	Status check	Alarm status	Availability Zone ↗	Public IP
<input type="checkbox"/>	worker-1_sb	i-078ed31a706f6f589	Running 🔍 🔍	t2.micro	2/2 checks passed	View alarms +	us-east-1a	ec2-54-89...
<input type="checkbox"/>	worker-2_sb	i-05da8301287468d59	Running 🔍 🔍	t2.micro	2/2 checks passed	View alarms +	us-east-1f	ec2-2-237...
<input type="checkbox"/>	kuber	i-097d2af538d0ca45a	Pending 🔍 🔍	t2.micro	-	View alarms +	us-east-1f	ec2-3-237...

- Thus Kuber named -instance gets created. Then click on Id of that instance then click on connect button you will see this:

[EC2](#) > [Instances](#) > [i-09dbca91fa3edcaea](#) > [Connect to instance](#)

Connect to instance [Info](#)

Connect to your instance i-09dbca91fa3edcaea (kuber) using any of these options

[EC2 Instance Connect](#) [Session Manager](#) [SSH client](#) [EC2 serial console](#)

Port 22 (SSH) is open to all IPv4 addresses
Port 22 (SSH) is currently open to all IPv4 addresses, indicated by 0.0.0.0/0 in the inbound rule in your [security group](#). For increased security, consider restricting access to only the EC2 Instance Connect service IP addresses for your Region: 18.206.107.24/29. [Learn more](#).

Instance ID
[i-09dbca91fa3edcaea](#) (kuber)

Connection Type

☒ Connect using EC2 Instance Connect
Connect using the EC2 Instance Connect browser-based client, with a public IPv4 address.

☐ Connect using EC2 Instance Connect Endpoint
Connect using the EC2 Instance Connect browser-based client, with a private IPv4 address and a VPC endpoint.

Public IPv4 address
[54.211.131.109](#)

Username
Enter the username defined in the AMI used to launch the instance. If you didn't define a custom username, use the default username, ec2-user.

Note: In most cases, the default username, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

[Cancel](#) [Connect](#)

- Then go into SSH client where you will get this command

Chmod 400 "keyname.pem"

ssh -i <keyname>.pem ubuntu@<public_ip_address> copy it and then connect it and run the following command for establishing connection.(I have entered this command on git bash where i entered in downloads where server.pem is stored then as the key is not accessible hence we need to change its mode using chmod 400 “key name.pem”. Then use the given command for making connections).

```
ADMIN@DESKTOP-LPV2RP5 MINGW64 ~
$ cd Downloads/
```

```
ADMIN@DESKTOP-LPV2RP5 MINGW64 ~/Downloads
$ chmod 400 "server.pem"
```

```
ADMIN@DESKTOP-LPV2RPS: MINGW64 ~/Downloads
$ ssh -i "server.pem" ec2-user@ec2-3-237-37-35.compute-1.amazonaws.com
The authenticity of host 'ec2-3-237-37-35.compute-1.amazonaws.com (3.237.37.35)'
can't be established.
ED25519 key fingerprint is SHA256:uhsLQC4WN+6i3np6iMFwW+vGF8aGkTrm89ryXQ4wTlg.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-3-237-37-35.compute-1.amazonaws.com' (ED25519) t
o the list of known hosts.
```

Amazon Linux 2023

<https://aws.amazon.com/linux/amazon-linux-2023>

```

Last login: Sat Sep 14 11:47:26 2024 from 18.206.107.28
[ec2-user@ip-172-31-65-181 ~]$

```

2. Installation of Docker

1. For installation of Docker into the machines run the following command: `sudo yum install docker -y`

```
[ec2-user@ip-172-31-65-181 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:08:49 ago on Sat Sep 14 11:42:25 2024.
Dependencies resolved.
=====
Package                Arch      Version                               Repository      Size
=====
Installing:
docker                 x86_64    25.0.6-1.amzn2023.0.2               amazonlinux     44 M
Installing dependencies:
containerd             x86_64    1.7.20-1.amzn2023.0.1               amazonlinux     35 M
iptables-libs          x86_64    1.8.8-3.amzn2023.0.2               amazonlinux     401 k
iptables-nft           x86_64    1.8.8-3.amzn2023.0.2               amazonlinux     183 k
libcgroup              x86_64    3.0-1.amzn2023.0.1                 amazonlinux     75 k
libnetfilter_conntrack x86_64    1.0.8-2.amzn2023.0.2               amazonlinux     58 k
libnftnl               x86_64    1.0.1-19.amzn2023.0.2              amazonlinux     30 k
libnftnl               x86_64    1.2.2-2.amzn2023.0.2              amazonlinux     84 k
pigz                   x86_64    2.5-1.amzn2023.0.3                 amazonlinux     83 k
runc                   x86_64    1.1.13-1.amzn2023.0.1              amazonlinux     3.2 M
=====
Transaction Summary
=====
Install 10 Packages

Total download size: 84 M
Installed size: 317 M
Downloading Packages:
(1/10): iptables-libs-1.8.8-3.amzn2023.0.2.x86_ 4.5 MB/s | 401 kB    00:00
(2/10): iptables-nft-1.8.8-3.amzn2023.0.2.x86_ 4.5 MB/s | 183 kB    00:00
(3/10): libcgroup-3.0-1.amzn2023.0.1.x86_64.rpm 3.9 MB/s | 75 kB    00:00
(4/10): libnetfilter_conntrack-1.0.8-2.amzn2023 2.6 MB/s | 58 kB    00:00
(5/10): libnftnl-1.0.1-19.amzn2023.0.2.x86_ 1.2 MB/s | 30 kB    00:00
(6/10): libnftnl-1.2.2-2.amzn2023.0.2.x86_64.rp 2.2 MB/s | 84 kB    00:00
(7/10): pigz-2.5-1.amzn2023.0.3.x86_64.rpm      2.8 MB/s | 83 kB    00:00
(8/10): runc-1.1.13-1.amzn2023.0.1.x86_64.rpm    30 MB/s | 3.2 MB   00:00
(9/10): containerd-1.7.20-1.amzn2023.0.1.x86_64 38 MB/s | 35 MB    00:00
(10/10): docker-25.0.6-1.amzn2023.0.2.x86_64.rp 34 MB/s | 44 MB    00:01
-----
Total                                     62 MB/s | 84 MB    00:01
-----
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.

Verifying : containerd-1.7.20-1.amzn2023.0.1.x86_64
Verifying : docker-25.0.6-1.amzn2023.0.2.x86_64
Verifying : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
Verifying : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
Verifying : libcgroup-3.0-1.amzn2023.0.1.x86_64
Verifying : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
Verifying : libnftnl-1.0.1-19.amzn2023.0.2.x86_64
Verifying : libnftnl-1.2.2-2.amzn2023.0.2.x86_64
Verifying : pigz-2.5-1.amzn2023.0.3.x86_64
Verifying : runc-1.1.13-1.amzn2023.0.1.x86_64

Installed:
containerd-1.7.20-1.amzn2023.0.1.x86_64      docker-25.0.6-1.amzn2023.0.2.x86_64      iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
libcgroup-3.0-1.amzn2023.0.1.x86_64          libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64      libnftnl-1.0.1-19.amzn2023.0.2.x86_64
pigz-2.5-1.amzn2023.0.3.x86_64              runc-1.1.13-1.amzn2023.0.1.x86_64

Complete!
```

- Then, configure cgroup in a daemon.json file by using following commands
`cd /etc/docker`
`cat <<EOF | sudo tee /etc/docker/daemon.json`

```
{
"exec-opts":
["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
```

```
}
EOF
```

```
[ec2-user@ip-172-31-65-181 ~]$ cd /etc/docker
[ec2-user@ip-172-31-65-181 docker]$ cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts":
  ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
{
  "exec-opts":
  ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
}
```

- Then after this run the following command to enable and start docker and also to load the daemon.json file.

```
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

```
[ec2-user@ip-172-31-81-216 docker]$ sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-81-216 docker]$ sudo systemctl daemon-reload
[ec2-user@ip-172-31-81-216 docker]$ sudo systemctl restart docker
```

- docker -v

```
[ec2-user@ip-172-31-65-181 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
```

3. Then Install Kubernetes with the following command.

- SELinux needs to be disabled before configuring kubelet thus run the following command
sudo setenforce 0

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

```
[ec2-user@ip-172-31-65-181 docker]$ sudo setenforce 0
[ec2-user@ip-172-31-65-181 docker]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

- Here We are adding kubernetes using the repository whose command is given below. cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]

```

name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF

```

```

[ec2-user@ip-172-31-65-181 docker]$ sudo setenforce 0
[ec2-user@ip-172-31-65-181 docker]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[ec2-user@ip-172-31-65-181 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/ enabled=1
gpgcheck=1 gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/ enabled=1
gpgcheck=1 gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[ec2-user@ip-172-31-65-181 docker]$

```

- After that Run following command to make the updation and also to install kubelet ,kubeadm, kubectl: `sudo yum update`

```

[ec2-user@ip-172-31-65-181 docker]$ sudo yum update
Invalid configuration value: gpgcheck=1 gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni in /etc/yum.repos.d/kubernetes.repo; invalid boolean value '1 gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni'
Error: Cannot find a valid baseurl for repo: kubernetes
Ignoring repositories: kubernetes
Last metadata expiration check: 0:19:14 ago on Sat Sep 14 11:42:25 2024.
Dependencies resolved.
Nothing to do.
Complete!

```

```

sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes

```

```

[ec2-user@ip-172-31-81-216 ~]$ sudo yum install -y kubelet kubeadm kubectl --dis
ableexcludes=kubernetes
Last metadata expiration check: 0:00:35 ago on Sat Sep 14 07:38:02 2024.
Dependencies resolved.

```

Package	Arch	Version	Repository	Size
Installing:				
kubeadm	x86_64	1.30.5-150500.1.1	kubernetes	10 M
kubectl	x86_64	1.30.5-150500.1.1	kubernetes	10 M
kubelet	x86_64	1.30.5-150500.1.1	kubernetes	17 M
Installing dependencies:				
conntrack-tools	x86_64	1.4.6-2.amzn2023.0.2	amazonlinux	208 k
cri-tools	x86_64	1.30.1-150500.1.1	kubernetes	8.6 M
kubernetes-cni	x86_64	1.4.0-150500.1.1	kubernetes	6.7 M
libnetfilter_cthelper	x86_64	1.0.0-21.amzn2023.0.2	amazonlinux	24 k
libnetfilter_cttimeout	x86_64	1.0.0-19.amzn2023.0.2	amazonlinux	24 k
libnetfilter_queue	x86_64	1.0.5-2.amzn2023.0.2	amazonlinux	30 k

Transaction Summary

```

Install 9 Packages

```

```

Total download size: 53 M
Installed size: 292 M

```

```

Downloading Packages:

```

```

(1/9): libnetfilter_cthelper-1.0.0-21.amzn2023. 436 kB/s | 24 kB 00:00
(2/9): libnetfilter_queue-1.0.5-2.amzn2023.0.2. 1.5 MB/s | 30 kB 00:00
(3/9): libnetfilter_cttimeout-1.0.0-19.amzn2023 309 kB/s | 24 kB 00:00
(4/9): conntrack-tools-1.4.6-2.amzn2023.0.2.x86 2.3 MB/s | 208 kB 00:00
(5/9): cri-tools-1.30.1-150500.1.1.x86_64.rpm 34 MB/s | 8.6 MB 00:00
(6/9): kubectl-1.30.5-150500.1.1.x86_64.rpm 21 MB/s | 10 MB 00:00
(7/9): kubeadm-1.30.5-150500.1.1.x86_64.rpm 17 MB/s | 10 MB 00:00
(8/9): kubelet-1.30.5-150500.1.1.x86_64.rpm 32 MB/s | 17 MB 00:00
(9/9): kubernetes-cni-1.4.0-150500.1.1.x86_64.r 21 MB/s | 6.7 MB 00:00

```

```

Verifying      : kubeadm-1.30.5-150500.1.1.x86_64 7/9
Verifying      : kubectl-1.30.5-150500.1.1.x86_64 7/9
Verifying      : kubelet-1.30.5-150500.1.1.x86_64 8/9
Verifying      : kubernetes-cni-1.4.0-150500.1.1.x86_64 9/9

```

```

Installed:

```

```

conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64
cri-tools-1.30.1-150500.1.1.x86_64
kubeadm-1.30.5-150500.1.1.x86_64
kubectl-1.30.5-150500.1.1.x86_64
kubelet-1.30.5-150500.1.1.x86_64
kubernetes-cni-1.4.0-150500.1.1.x86_64
libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64

```

```

Complete!

```

```

[ec2-user@ip-172-31-81-216 ~]$

```

- After installing Kubernetes, we need to configure internet options to allow bridging.
 - sudo swapoff -a
 - echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
 - sudo sysctl -p

```

[ec2-user@ip-172-31-81-216 ~]$ sudo swapoff -a
[ec2-user@ip-172-31-81-216 ~]$ echo "net.bridge.bridge-nf-call-iptables=1" | sud
o tee -a /etc/sysctl.conf
net.bridge.bridge-nf-call-iptables=1
[ec2-user@ip-172-31-81-216 ~]$ sudo sysctl -p
net.bridge.brdaae-nf-call-iptables = 1

```

4. Initialize the Kubecluster

```

sudo kubeadm init --pod-network-cidr=10.244.0.0/16

```

```

[ec2-user@ip-172-31-80-126 docker]$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
I0913 10:32:44.629146 26680 version.go:256] remote version is much newer: v1.31.0; falling back to: stable-1.30
[init] Using Kubernetes version: v1.30.4
[preflight] Running pre-flight checks

```


Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.26.174:6443 --token pv0yyi.xh1lqhclfjr50pt8 \
--discovery-token-ca-cert-hash sha256:8293b2f6d29de466bd859007f5adbcdb3a
ecb0c446ba09033d32a5846b3d434f
```

- copy the token and save for future use .

```
kubeadm join 172.31.26.174:6443 --token pv0yyi.xh1lqhclfjr50pt8
--discovery-token-ca-cert-hash
sha256:8293b2f6d29de466bd859007f5adbcdb3a
ecb0c446ba09033d32a5846b3d434f
```

- Copy the mkdir and chown commands from the top and execute

```
them mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```



```
[ec2-user@ip-172-31-80-126 docker]$ %C
[ec2-user@ip-172-31-80-126 docker]$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Then, add a common networking plugin called flannel as mentioned in the code.

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

```
[ec2-user@ip-172-31-26-174 docker]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
```

5. Now that the cluster is up and running, we can deploy our nginx server on this cluster. Apply deployment using this following command:

```
kubectl apply -f
```

```
https://k8s.io/examples/pods/simple-pod.yaml
```

```
[ec2-user@ip-172-31-26-174 docker]$ kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml
pod/nginx created
```

Then use kubectl get nodes to check whether the pod gets created or not.

```
[ec2-user@ip-172-31-26-174 docker]$ kubectl get pods
NAME    READY   STATUS    RESTARTS   AGE
nginx   0/1     Pending   0           12s
```

To convert state from pending to running use following command:

kubectl describe pod nginx This command will help to describe the pods it gives reason for failure as it shows the untolerated taints which need to be untainted.

```
[ec2-user@ip-172-31-26-174 docker]$ kubectl describe pod nginx
Name:          nginx
Namespace:     default
Priority:       0
Service Account: default
Node:          <none>
Labels:        <none>
Annotations:   <none>
Status:        Pending
IP:            <none>
IPs:           <none>
Containers:
  nginx:
    Image:          nginx:1.14.2
    Port:           80/TCP
    Host Port:      0/TCP
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-k4lj6 (ro)
```

```
    /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-k4lj6 (ro)
```

```
Conditions:
  Type             Status
  PodScheduled     False
Volumes:
  kube-api-access-k4lj6:
    Type:          Projected (a volume that contains injected data from m
multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:        kube-root-ca.crt
    ConfigMapOptional:    <nil>
    DownwardAPI:          true
QoS Class:           BestEffort
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 3
00s
                    node.kubernetes.io/unreachable:NoExecute op=Exists for
300s
```

```
Events:
  Type    Reason             Age   From                    Message
  ----    -
  Warning FailedScheduling   7s    default-scheduler      0/1 nodes are available: 1 no
de(s) had untolerated taint {node-role.kubernetes.io/control-plane: }.
preemption: 0/1 nodes are available: 1 Preemption is not helpful for scheduling.
```

```
[ec2-user@ip-172-31-26-174 ~]$ kubectl taint nodes --all node-role.kubernetes.io
/control-plane-
node/ip-172-31-26-174.ec2.internal untainted
```

6. Now check pod status is is running

```
[ec2-user@ip-172-31-26-174 ~]$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   1 (6s ago)  90s
```

7. Lastly, mention the port you want to host. Here i have used localhost 8081 then check it.

```
kubectl port-forward nginx 8081:80
```

```
[ec2-user@ip-172-31-26-174 ~]$ kubectl port-forward nginx 8081:80
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

8. Verify your deployment

Open up a new terminal and ssh to your EC2 instance.

Then, use this curl command to check if the Nginx server is running.

```
curl --head http://127.0.0.1:8080
```

If the response is 200 OK and you can see the Nginx server name, your deployment was successful. We have successfully deployed our Nginx server on our EC2 instance.

Conclusion:

First, I successfully launched an AWS EC2 instance running Amazon Linux. After that, I installed Docker and Kubernetes on the instance. Following the installation, I initialized the Kubernetes cluster, which provided me with a token, along with chown and mkdir commands. I then executed both the mkdir and chown commands successfully. Next, I installed the Flannel networking plugin without any issues. Initially, there was an error while deploying Nginx, but after correcting it, I successfully deployed Nginx using a simple-pod.yml file. I confirmed its deployment with the get pods command and hosted it locally on <http://localhost:8081>, which worked as expected.